

Forecasting Hourly Household Electricity Consumption: A Leakage-Safe Comparison of Classical, Tree-Based, and Sequence Models

Thierry Ishimwe | Linda Carla Zorzoli | Mariavittoria Giurato | Cesar Dushimimana
LUISS Guido Carli · Artificial Intelligence Techniques · Industry partner: Enel Global ICT · May 2026

Executive Summary

This report compares four model families — SARIMA, XGBoost, LSTM, and GRU — on the task of forecasting hourly residential electricity consumption (UCI *Individual Household Electric Power Consumption* dataset, Sceaux, France, 2006–2010, 2,075,259 minute rows resampled to 34,000 hourly observations). The pipeline is leakage-safe by construction: every feature is registered with a leakage-safety flag and a written rationale, same-timestamp components of the target (*Voltage*, *Global_intensity*, sub-metering) are mechanically excluded, and short-gap imputation is past-only so no future-derived value can flow into a training row. Cross-fold mean MAE for the winning model (XGBoost) is **0.336 kW**.

Seven findings shape the conclusion:

- **XGBoost wins the leaderboard at MAE = 0.336**. Its lead is statistically significant (Diebold–Mariano, $\alpha = 0.05$) only against SARIMA ($p = 0.022$). Against LSTM ($p = 0.33$) and GRU ($p = 0.42$) the difference is within noise — the top three are a statistical tie.
- **The four real model families cluster within ~3% MAE** (XGBoost 0.336, GRU 0.339, LSTM 0.341, SARIMA 0.345). Architecture choice within this band is a near-tie; deployment should be driven by non-MAE criteria.
- **Peak-hour MAE (18–21h) is ~2.5× trough-hour MAE (03–06h)** for every model. XGBoost is best at quiet hours; GRU is slightly better at peak hours — a real operational consideration since peak forecast errors drive imbalance penalties.
- **October is the hardest validation month** (heating-ramp, MAE ≈ 0.44); **August is easiest** (the French *grandes vacances* — empty household, flat consumption, MAE ≈ 0.22). The pattern is structural, not a single-fold artifact.
- **Stacking ensembles do not break the 3% cluster**. A ridge meta-learner over the four base models edges past every individual model on the held-out fold (0.368 vs 0.368 best individual); base-model errors are too correlated for simple linear stacking to help.
- **Weather augmentation validates the "input information, not architecture" hypothesis**. Augmenting XGBoost with four Paris–Montsouris weather columns (temperature, humidity, wind speed, 24h temperature lag) reduces cross-fold MAE from 0.336 to 0.334 (−0.77%), with the improvement concentrated on the hardest heating-season folds.
- **The 1-step-ahead hourly horizon is a methodology benchmark, not the operational target**. Production day-ahead forecasts at Enel use 24-step horizons made from a fixed daily cutoff. The reported MAE values are therefore a lower bound on operationally-realistic difficulty; the protocol switch is recorded as future work.

1. Problem and Dataset

We forecast hourly *Global_active_power* (kW) for a single household in Sceaux, Paris area, using the UCI *Individual Household Electric Power Consumption* dataset: 2,075,259 minute-level observations from 16 December 2006 to 26 November 2010 (47 months, 7 measured variables). The source file is pinned by SHA-256 in `conf/base.yaml` and is verified on download. The 1-minute series is resampled to hourly mean (34,000 hourly rows), aligning with the 1-hour bidding blocks used in European electricity markets.

Missingness is 1.25% (25,979 minute rows across 71 distinct gap segments), bimodally distributed: 54 short gaps (≤ 3 min, 72 minutes total) and 17 long outages (> 3 min, 25,907 minutes total). The longest single outage spans 5 days in August 2010, coinciding with the French summer-vacation period. Short gaps are forward-filled from the last past observation (past-only, no future leakage); long outages are preserved as NaN with explicit *is_outage_gap* indicators. After hourly resampling, 28 hourly rows (0.08%) retain the outage flag.

2. Methodology

2.1 Leakage-Safe Feature Engineering

Every feature is registered in `features/registry.py` with a leakage-safety flag and a written rationale. The pipeline refuses to emit any column marked unsafe when `strict_leakage_safe: true`. Same-timestamp UCI columns that are physical components of the target (*Voltage*, *Global_intensity*, sub-metering 1/2/3, *Global_reactive_power*) are registered as unsafe and blocked from entering the feature matrix — they would let the model recover the target through $P = V \cdot I$ rather than forecast it. Short-gap imputation uses past-only forward-fill in `preprocessing/missing.py`, so no future observation can leak into a training row through the cleaned series. Two regression tests guard the policy — `tests/integration/test_leakage.py` (registry contract) and `tests/unit/test_missing_value_policy.py` (past-only fill).

The leakage-safe feature matrix has 31 columns across five families: **target lags** at 1, 24, and 168 hours (one hour, one day, one week back); **rolling statistics** (mean, std, min, max) over 3, 24, and 168-hour windows of the lag-1-shifted target — the *shift-then-roll* order guarantees the window at time t covers strictly past values; **calendar features** (hour, day-of-week, month, *is_weekend*, *is_french_holiday*); **cyclical encodings** (*sin/cos* of hour, dow, month) so trees can handle the wraparound at midnight and end-of-week; and **outage indicators**.

2.2 Evaluation Protocol

Rolling-origin cross-validation, 6 folds, each fold trains on 365 days and validates on the following 30 days. The folds slide forward in 30-day increments; validation never overlaps training and earlier folds never see later data. Each model is fit from scratch per fold. The protocol is **1-step-ahead at hourly resolution**: at validation hour t , the model predicts $y[t]$ using information available through $t-1$. Metrics computed per fold: MAE (primary), RMSE, WAPE, sMAPE, MASE (with seasonal period 24 for the daily-naive denominator).

True 24-step-ahead day-ahead forecasting — the horizon that matches Enel's market-bid block — is listed as future work; the reported MAE values are therefore a lower bound on the operationally-realistic difficulty.

3. Model Lineup

The lineup spans four model families. Each was selected to test a distinct hypothesis about which class of structure dominates the signal.

Model	Family	Role
naive_lag_{1h, 24h, 168h}	Persistence baselines	Floors for daily / weekly seasonality
SARIMA(1,1,1)(1,1,1,24)	Classical statistical	Satisfies brief's ≥ 1 classical requirement
LSTM (168h lookback, 64 units)	Deep-learning sequence	Satisfies brief's ≥ 1 DL requirement
GRU (168h lookback, 64 units)	Deep-learning sequence	Within-family DL comparison
XGBoost (depth=6, lr=0.05, n=500)	Gradient-boosted trees	Tabular model with 31 engineered features
Stacking (ridge $\alpha=1.0$)	Ensemble (bonus)	Tests base-model error decorrelation

4. Results

4.1 Leaderboard




Figure 1. Cross-fold mean MAE per model with $\pm 1\sigma$ error bars across the six rolling-origin folds. Red bars are naive persistence baselines; blue bars are the real model families. Dashed line marks the lag-1h persistence floor (0.380).

Model	MAE	RMSE	MASE
XGBoost	0.337 $\pm$ 0.068	0.468 $\pm$ 0.093	0.573 $\pm$ 0.119
GRU	0.339 $\pm$ 0.086	0.475 $\pm$ 0.103	0.578 $\pm$ 0.149
LSTM	0.341 $\pm$ 0.065	0.470 $\pm$ 0.091	0.580 $\pm$ 0.114
SARIMA	0.345 $\pm$ 0.079	0.490 $\pm$ 0.097	0.588 $\pm$ 0.137
Naive lag-1h	0.380 $\pm$ 0.107	0.563 $\pm$ 0.135	0.647 $\pm$ 0.186
Naive lag-24h	0.487 $\pm$ 0.153	0.715 $\pm$ 0.180	0.828 $\pm$ 0.264
Naive lag-168h	0.546 $\pm$ 0.131	0.777 $\pm$ 0.170	0.929 $\pm$ 0.230

Table 1. Cross-fold mean $\pm$ std for each model. Lower is better. XGBoost wins on MAE; all four real models cluster within 3% of each other.

4.2 Statistical Significance — Diebold–Mariano Test

Pairwise Diebold–Mariano tests on absolute-error loss between the four real models, using the Harvey–Leybourne–Newbold small-sample correction and Newey–West HAC variance at lag 0 (appropriate for 1-step-ahead forecasts). Two-sided p-values:

	XGBoost	SARIMA	LSTM	GRU
XGBOOST	—	0.022	0.331	0.415
SARIMA	0.022	—	0.108	0.073
LSTM	0.331	0.108	—	0.747
GRU	0.415	0.073	0.747	—

Table 2. DM two-sided p-values, $n \approx 4,100$ paired predictions. Only XGBoost vs SARIMA ($p = 0.022$) clears the $\alpha = 0.05$ bar. XGBoost vs LSTM ($p = 0.33$) and XGBoost vs GRU ($p = 0.42$) are not statistically distinguishable.

4.3 Per-Segment Performance




Figure 2. MAE decomposed by hour-of-day, day-of-week, and month for each of the four real models. Aggregate MAE hides operationally important structure.

Hour of day. Every model exhibits the same shape: easy at 03–06h (MAE ≈ 0.17), hard at 18–21h (MAE ≈ 0.45). XGBoost wins quiet hours; GRU is slightly better at peak hours. **Day of week.** Weekend MAE is $\sim 5\%$ higher than weekday — modest. **Month.** August is by far the easiest validation month (MAE 0.21–0.25, the empty-household *grandes vacances*); October is the hardest (0.42–0.44, the regulated *période de chauffe* heating ramp).

4.4 Weather Data as Exogenous Regressor

To test the "input information, not architecture" hypothesis empirically, we retrained XGBoost on the feature matrix augmented with four hourly weather columns from station Paris–Montsouris (5.2 km from Sceaux, via the *meteostat* client): temperature ($^{\circ}\text{C}$), relative humidity (%), wind speed (km/h), and a strictly past 24-hour-lagged temperature. Per-fold result:

Fold	Baseline MAE	+Weather MAE	Δ MAE	Δ %
1	0.3351	0.3359	+0.0008	+0.24%
2	0.3324	0.3269	-0.0055	-1.65%
3	0.2147	0.2122	-0.0025	-1.16%
4	0.3352	0.3320	-0.0032	-0.95%
5	0.4024	0.3999	-0.0025	-0.62%
6	0.3994	0.3967	-0.0027	-0.68%
Mean	0.3365	0.3339	-0.0026	-0.77%

Table 3. XGBoost with vs without four weather features. The mean MAE delta is small (-0.77%) but signed in the predicted direction and concentrated on the harder folds: the late-winter / spring / October-heating-ramp windows show the larger gains, while the easy summer fold (1) gains little — consistent with weather mattering when consumption is temperature-driven, not when the household is flat-line.

5. External Context

Several empirical patterns map onto documented French and European context. The longest outage (17–22 August 2010) and the "August is easiest" finding both reflect the institutionalised summer vacation that INSEE and RTE document accounts for a national consumption drop of 15–20%. The January 2010 outage (12–14 January) coincides with a European cold wave during which RTE's *Bilan Électrique 2010* records French national load exceeding 96 GW for the first time, driven by the country's $\sim 30\%$ electric-heating penetration. October's role as the hardest validation month is explained by the regulated *période de chauffe* beginning 15 October each year, after which RTE documents typical residential consumption rising 40–60% between mid-October and mid-November.

6. Discussion

Architecture matters less than engineering protocol. Four model families — classical SARIMA, tree-based XGBoost, sequence DL LSTM and GRU — cluster within 3% MAE, and the within-cluster differences are statistically indistinguishable. This is consistent with M-competition findings that on data-rich forecasting problems, feature engineering and evaluation discipline dominate architecture choice. The stacking-ensemble result corroborates this: base-model errors are correlated enough that linear combination does not help — ruling out "more model variety" as the path forward.

The "best" model depends on the operational metric. XGBoost wins overall MAE but GRU wins peak-hour (18–21h) MAE. For a grid operator whose cost function weights peak-hour accuracy heavily — a real concern since peak forecast errors drive imbalance penalties — GRU is the more defensible deployment choice despite XGBoost's leaderboard lead. This nuance is hidden by an aggregate MAE.

GRU is competitive with LSTM at lower complexity. Across all six folds, GRU's MAE (0.339) is essentially tied with LSTM's (0.341). The extra output gate in LSTM is unused capacity on this dataset; GRU is a reasonable default when adding a sequence-DL model to a forecasting pipeline.

Weather augmentation: small but real, where it matters. A -0.77% MAE gain on XGBoost from four free Paris–Montsouris weather columns is modest, but signed in the predicted direction and concentrated on heating-season folds. It validates the direction of the priority claim — better input information is closer to the leaderboard ceiling than another architecture would bring us — without overclaiming a transformative gain.

7. Future Work

- **Extend the weather-augmented matrix to SARIMA, LSTM, and GRU.** The augmented feature matrix is already persisted at `data/processed/features_with_weather.parquet`. The missing work is ~ 70 min of model retraining. Expected to confirm whether the -0.77% XGBoost gain is single-model-specific or a team-wide improvement.
- **Implement 24-step-ahead day-ahead forecasting.** Current results are 1-step-ahead at hourly resolution. Production day-ahead forecasts at Enel use 24-step horizons made once per day from a fixed daily cutoff; the reported MAE values are a lower bound on operationally-realistic difficulty.

8. Reproducibility

Every result above is reproducible from a fresh clone of github.com/ThierryIshimwe/Energy-Forecasting. The deliverable notebook (`notebooks/00_main_deliverable.ipynb`) loads tracked artifacts only and ships with outputs pre-rendered, so an evaluator needs three commands: `clone, pip install -r requirements.txt && pip install -e .`, open the notebook. No raw-data download is needed for the deliverable view; the EDA summary numbers and seasonality plot data are precomputed and tracked.

Pipeline guarantees: Python 3.11 pinned in `.python-version`; all dependencies version-pinned in `requirements.txt`; global RNG seed 42 in `conf/base.yaml`; source dataset SHA-256 verified at download; 64+ unit + integration tests, including the leakage-registry test and the past-only-imputation test, run on every commit. Re-training every model from raw data takes ~ 90 minutes and is documented in *README.md*.